

# Fundamentals of ML

Mehdi Rostami

University of Toronto

Bank of Montreal

September 4, 2023

# Regression function

- Find  $a(X)$  that minimizes the risk  $\mathbb{E}(Y - a(X)|X)^2$ ?
- $a(X) = \mathbb{E}(Y|X)$
- This is called "regression function" that in ML we want to estimate as it minimizes the risk
- When  $Y$  is binary,  $\mathbb{E}(Y|X) = P(Y = 1|X)$ 
  - What is the expected value of Bernoulli distribution?
  - $Z \sim \text{bernoulli}(p)$  then  $E(Z) = p$
  - $Z|W \sim \text{bernoulli}(P(Z = 1|W))$  then  $E(Z|W) = P(Z = 1|W)$

# What is Learning?

- PAC learning: Probably approximately correct
- Roughly speaking, a class of algorithms (like the class of all neural nets) is called PAC learnable, if, for any dataset, and any  $\epsilon$  and  $\delta$  (small positive numbers), there's an algorithm in that class so that the error (MSE) is smaller than  $\epsilon$  with probability at least  $1 - \delta$ .

# Model vs algorithm

- Different ML algorithms try to estimate this function using the input data using some universal function approximator like polynomial function (Taylor expansion), tree-based algorithms, splines, neural networks, etc.
- Note that I'm saying "algorithm", not models. People use the word "model" but that's technically wrong. These are algorithms or function approximators; we minimize the error between the observed and predicted values so they pick the best algorithm among the class of all such algorithms. We are just minimizing a loss (that's not coming from a distribution assumed for the outcome-input relationship); so we're not modeling the reality.
- In addition, in practice, we use many "algorithms" and pick the best one; this is not the right way of modeling things. To model the reality, you assume some form for the reality and use data to show that, and that's all.
- All ML algorithms are non parametric. Sure, we assume that the regression function belongs to  $\{P_\theta|\theta\}$ , but there's no restriction on  $\theta$ .

# No-Free-Lunch Theorem

- For every algorithm there exists a distribution (or data) on which it fails. In other words, there's no algorithm that's the best for all tasks.
- Pick an algorithm that you claim to be best for all datasets. I'll generate a dataset that some other algorithm will outperform your candidate.
- Summary: There is no universal learner: Every learner has to be specified to some task, and use some prior knowledge about that task (which class to choose), in order to succeed.

# Linear Regression Model

- Goal: Calculate the first-order approximation of the Taylor expansion of  $\mathbb{E}(Y|X)$
- Assumptions: Nothing! No linearity, normality or conditional IID is needed.
- For prediction, we don't need distributional assumptions; With the right interpretation of coefficients, linearity is also not needed
- Goal: Estimate the parameters

$$\mathbb{E}[Y|X] = X\beta = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p \quad (1)$$

- Once somehow estimated, prediction for a new observation  $X = x$ :

$$\hat{\mathbb{E}}[Y|X = x] = x\hat{\beta} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p \quad (2)$$

- With either assuming that the data is normal or just using least square method, we have this objective function:  $\min_{\beta} \mathbb{E}(Y - X\beta)^2$  which for a finite sample amounts to  $\min_{\beta} \sum_{i=1}^n (y_i - X_i\beta)^2$
- A closed form solution is

$$\hat{\beta} = (X^T X)^{-1} X^T y$$


- The prediction for a new  $x$  is  $x\hat{\beta} = x(x^T x)^{-1} x^T y$
- If  $p < n$  and  $X^T X$  is not singular, then  $\hat{\beta}$  is a consistent and unbiased estimator of  $\beta$ , and if normality of  $y$  is assumed, it's asymptotically normal too (good for confidence interval calculation).

# Bias-Variance Tradeoff

- Let  $f(X)$  be the best approximator in a class of algorithms (e.g. best decision-tree/linear regression) of  $Y$  and  $\hat{f}(X)$  be it's data-driven estimator,

$$\mathbb{E} [Y - \hat{f}(X)]^2 = \mathbb{E} [f(X) - \hat{f}(X)]^2 + \mathbb{E} [Y - f(X)]^2 =$$
$$\left\{ \mathbb{E} [\hat{f}(X) - \mathbb{E}[\hat{f}(X)]]^2 + (\mathbb{E} [f(X) - \mathbb{E}[\hat{f}(X)]] )^2 \right\} + \left\{ \mathbb{E} [Y - \mathbb{E}[Y|X]]^2 + \mathbb{E} [\mathbb{E}[Y|X] - f(X)]^2 \right\}$$

- Error = Estimation Error + Approximation Error =  
{Estimator Variance + Estimator Bias} + {Measurement error + Class-specific error}
- Example:**
  - $f(x)$  is a linear regression and  $\mathbb{E} [Y - \hat{f}(X)]^2 = 4 (1+0+\mathbf{1}+2)$
  - $f(x)$  is a decision tree and  $\mathbb{E} [Y - \hat{f}(X)]^2 = 4 (1.5+.5+\mathbf{1}+1)$





# Interpretation of Bis-Variance components

- **Estimation Error** is summation of the variance of the data-driven estimator and its bias from the best un-realistic estimator in the chosen class of algorithms.
  - ① Depends on both data and the chosen class
  - ② The bias component can be reduced by choosing a different class of algorithms. Unchanged by sample size!
  - ③ The variance component can be reduced if, for a chosen estimator, you force it to less focus on data.
- **Approximation Error** is the summation of measurement error in the data-generating process and the error of the chosen class of algorithms
  - ① Unchanged by data or its size.
  - ② The measurement error component is due to an error in  $Y$  measurements; can't be reduced.
  - ③ The class-specific error is induced by the class of algorithms we pick. Can't reduce it after you choose a class! Can only reduce it by choosing a more complex class of algorithms.

# Predictions variance in linear regression

- For linear regression, the in-sample error and  $\mathbb{E}(X^T \beta - X^T \hat{\beta})^2 = c \frac{p}{n}$  out-of-sample error is  $\mathbb{E}(x^T \beta - x^T \hat{\beta})^2 = c \frac{p}{n-p-1}$
- When  $p \ll n$  things are well.
- What happens if  $p < n$  but  $p \rightarrow n$ ? The model doesn't generalize well.
- How about  $p \geq n$ ? Need hard-feature selection
- By increasing # of features in the model, complexity increases  $\rightarrow$  bias goes down, variance goes up
- The above shows the dimensionality effect or curse of dimensionality in linear regression. It's valid in other algorithms too.
- Curse of dimensionality: Error increases by # of features.

# Bias-variance of linear regression

- For linear regression predictions:

$$\begin{aligned}\mathbb{E} [Y - X\hat{\beta}]^2 &= \mathbb{E} [X\hat{\beta} - \mathbb{E}[X\hat{\beta}]]^2 + \left( \mathbb{E} [X\beta - \mathbb{E}[X\hat{\beta}]] \right)^2 + \mathbb{E} [Y - X\beta]^2 = \\ &\mathbb{E} [X\hat{\beta} - X\beta]^2 + \mathbb{E} [Y - X\beta]^2\end{aligned}$$

- Since the linear regression parameters are unbiased, the second term is zero! So, in linear regression, we have no bias!
- Why linear regression is not the best:
  - If  $p \ll n$  then the approximation error is very high
  - If  $p$  is large, not only the approximation error is high (although improved compared to low  $p$ ), but the variance is high!

# Bias-variance tradeoff - con't

- To minimize risk, there's a tradeoff, called the bias-variance tradeoff.
- Choosing the class of algorithms to be a very rich class decreases the approximation error but at the same time might increase the estimation error, as a rich class might lead to overfitting.
- Choosing the class of algorithms to be a very small set reduces the estimation error but might increase the approximation error or, in other words, might lead to underfitting.
- To have intuition about why there's a tradeoff, consider the extreme. The lowest bias can be achieved by a perfect fit of the data (zero error), but it doesn't generalize to unseen data at all, meaning high variance. On the other hand, the lowest variance is a constant, but it will have huge errors and is badly biased.
- Bias and variance cannot be directly observed, because  $\hat{\theta}$  calculated on the data is a number that we cannot find its expectation accurately. Bootstrap can help estimate them
- The training MSE gives an intuition about the bias, and the validation MSE gives an intuition about the variance.

# Regularization

- A solution to avoid an increase in variance is regularization
- Any modification to an algorithm that decreases the variance while increase the bias (class-specific error) can be regarded as a regularization method
- We sacrifice lower bias to get lower variance; In fact, the model complexity goes down by regularization; increasing one component of the error and decreasing another
- One way to impose regularization in loss minimization problems is to shrink the model parameters while estimating them during the optimization.
- Search  $\hat{\beta}$ 's that minimize the least squares in a constrained subspace

$$\min_{\beta} \mathbb{E}(Y - X\beta)^2, \quad \text{subject to}$$
$$h(\beta) < \frac{c}{\lambda}$$

- ATT conditions in convex optimization  $\min_{\beta} \mathbb{E}(Y - X\beta)^2 + \lambda h(\beta)$
- This is not the only way of imposing regularization to ML algorithms. There are a few more that will be reviewed later. However, the above way of regularizing help study the properties.

# Examples




- ATT conditions in convex optimization  $\min_{\beta} \mathbb{E}(Y - X\beta)^2 + \lambda h(\beta)$
- **Ridge** ( $L_2$ ):  $h(\beta) = \|\beta\|^2 \rightarrow$  Shrinkage and unique solution (the latter was the trigger for discovering the idea of regularization)
- **LASSO** ( $L_1$ ):  $h(\beta) = \|\beta\| \rightarrow$  Sparsity and shrinkage. Could generate multiple solutions.
- **Elastic net**:  $h(\beta) = \|\beta\| + \|\beta\|^2 \rightarrow$  Inherits both properties

# Visual interpretation

- For example:  $\min_{\beta} \mathbb{E}(Y - X\beta)^2 + \lambda \|\beta\|$  or

$$\min_{\beta} \mathbb{E}(Y - X\beta)^2, \quad \text{subject to}$$

$$\|\beta\| < \frac{c}{\lambda}$$



- Very effective method in practice as compared to other linear models
- The best first method to try for feature selection
- Selective Inference can give uncertainty measures (p-values) for the estimators of the parameters.
- $\lambda$  is a measure of complexity; very large values of  $\lambda$  zero out all coefficients, while smaller values could cause overfitting. Requires hyperparameter tuning to select  $\lambda$ .
- Larger the regularization parameter  $\lambda$ , tighter the constraint optimization and thus more shrinkage.
- Extreme cases  $\lambda = 0$  we impose no regularization and  $\lambda = \infty$  the algorithm learns nothing



# Bayesian interpretation

- Assuming **conditional normality** for  $y$ , the log posterior becomes the regulation objective - because posterior is the multiplication of likelihood and prior.
- The normal prior for  $\beta \rightarrow$  Ridge
- The Laplace prior for  $\beta \rightarrow$  LASSO
- How? Because posterior  $\propto$  likelihood  $\times$  prior:

$$P(\beta|Y, X) \propto e^{-\sum_{i=1}^n (y_i - x\beta)^2} \times e^{-\alpha \|\beta\|^2},$$

$$P(\beta|Y, X) \propto e^{-\sum_{i=1}^n (y_i - x\beta)^2} \times e^{-\alpha \|\beta\|}.$$

- minimizing log posterior is the same as minimizing the regularized loss seen before.
- Intuition: In regularization we have less focus on the data  $\rightarrow$  We have less variance (remember the statement from a few pages back)

# Properties of regularized estimators

- Helps bias-variance tradeoff for predicted values; overall better fit
- First consequence  $\rightarrow$  the regularized estimator is biased.
  - That is  $\mathbb{E}[\hat{\beta}_{regularized}] \neq \beta$ . Easy to see in the case of Ridge.
- Second consequence  $\rightarrow$  the regularized estimator is less variable.
  - That is  $Var(\hat{\beta}_{regularized}) \leq \hat{\beta}_{OLS}$
  - Easy to see for Ridge when orthogonal covariates,  $\hat{\beta}_{ridge} = \frac{n}{n+\lambda} \hat{\beta}_{OLS}$
  - Similar result in non-orthogonal case
  - Or with the change in the data, the regularized estimator varies less.
  - Intuition for lower var: It's as if we focus less on the data by introducing an informative prior.
  - Variance of predictions is low:  $X\hat{\beta}_{regularized}$ ; the lower the values of  $\beta$ , the less sensitive to the variations (or errors/deviations) from  $X$ .
  - Stability of algorithm when  $L_2$  regularization is used is  $\propto \frac{1}{\lambda n}$
  - Stable algorithms don't overfit
- How can we see in practice/simulations the variance of other regularized estimators?

# Other types of regularization

A regularization technique doesn't necessarily need to add a penalty to the loss function! Anything that decreases the model complexity (and the variance) can be regarded as regularization.

- Early stopping
- Batch learning (randomly) imposes some regularization (Like in NNs)
- Inputting a random subset of features into the model (like a random forest)
- Ensembling
- Dropout
- ?

Drawback of informal regularization: Hard to study the theoretical properties.

Think about how these regularizations decrease focus on the data that help reduce variance of the estimator.